

nanoGenRec: An Agentic Research Framework for Landing Semantic-ID Generative Recommendation

Ziqing Ye
yeziqing986@gmail.com

June 12, 2026

Abstract

AI agents can now read papers, write code, and launch experiments, but research progress still depends on a harder layer: maintaining a reproducible operating system around hypotheses, configs, long-running jobs, evaluation contracts, failed runs, and human decisions. We present NANOGENREC, an open-source agentic research framework that makes this layer explicit. The framework provides an inbox/outbox protocol for human-agent collaboration, paper-note memory, YAML experiment configuration, duplicate-run detection, runtime estimation, queue-based execution, full-recall evaluation checks, decision records, and durable experiment logs. Semantic-ID generative recommendation is used as the stress test because it is a realistic multi-stage research problem rather than a toy workflow: item representation, tokenization, behavior-sequence construction, next-token training, reward alignment, constrained decoding, and full-recall evaluation must all agree before any result is meaningful. The project is distilled from an anonymized production-scale experiment lineage: 51 recorded experiments, behavior sweeps up to 7.85M users and 299.0M raw interactions, roughly 445M effective SID tokens after truncation, million-item SID caches, and experiments run in a real 4×NVIDIA H200 GPU environment. Within this agent-managed workspace, the GR stack serves as a validation target and produces traceable evidence of model scaling, non-monotonic behavior-window effects under temporal drift, tokenizer-capacity tradeoffs, SID-constrained retrieval, and post-training alignment failures caused by rollout distribution mismatch. The released artifact removes private data and deployment details while preserving the research machinery: configs, logs, validation checks, public notebooks, and reproducible code paths. A checked-in MovieLens 1M Colab-T4 path verifies that the public code can execute the same Qwen embedding → SID construction → NTP training → reward-alignment → full-evaluation loop without private data. We position NANOGENREC primarily as an agentic research workspace and landing record: a way for researchers and agents to turn complex recommender-system ideas into auditable, recoverable, executable research programs.

1 Introduction

Generative Recommendation (GR) is moving recommender systems from multi-stage retrieve-and-rank pipelines toward autoregressive generation over item tokens. Industrial reports such as OneRec [5], OneRec-V2 [6], OneMall [4], GenRec [7], OneLive [2], and GR4AD [3] show that Semantic-ID generation, model scaling, and preference alignment can be viable at production scale.

The public literature now contains strong system diagrams and deployment evidence, yet there is still a practical gap for researchers who want to turn those ideas into a running system. A working GR system needs more than a model definition. It needs a tokenizer that maps items into discrete semantic tokens, a behavior preprocessing path that produces autoregressive training sequences,

a model and training loop, a constrained generation procedure that maps generated tokens back to items, a full-recall evaluation protocol, and experiment infrastructure that can distinguish real improvements from duplicated baselines, invalid feature plumbing, or inline-evaluation artifacts. In practice, the system either works end to end or the individual components are hard to interpret.

We introduce NANOGENREC, an open-source agentic research framework and empirical record for landing Semantic-ID-based generative recommendation. The project is built from production-grounded experiments, with private data and deployment-specific details removed. Its purpose is not to claim a new recommendation algorithm. Instead, it shows how a research agent and human collaborator can assemble, execute, evaluate, debug, and iterate on a GR stack in a reproducible workspace. The accompanying experiment lineage records both positive results and failure modes, making the framework useful as a reference implementation, an operational checklist, and a durable research memory.

The claim boundary is explicit. NANOGENREC does not claim online production gains, a new tokenization algorithm, or a new preference-optimization objective. Its contribution is a landed research framework: the agent protocol, code paths, experiment contracts, evaluation protocol, and logs needed to make Semantic-ID GR run as a real retrieval system and to reproduce a set of useful GR behaviors under that framework.

This report makes five contributions:

1. **An agentic research framework.** We release an inbox/outbox protocol, paper-note memory, status tracking, decision records, runtime estimation, YAML experiment configuration, duplicate-run checks, and queue-based long jobs for recommendation research.
2. **A nontrivial GR validation target.** We instantiate the framework on an end-to-end Semantic-ID GR workflow covering tokenization, NTP training, reward alignment, SID-constrained beam search, and full-recall evaluation.
3. **Agent-managed empirical lineage.** Within the same framework, we preserve model scaling, tokenizer-capacity effects, full-recall baselines, and post-training alignment dynamics as traceable experiment records rather than isolated result tables.
4. **Production-grounded scale and public proof.** We report a seven-point NTP scaling sweep, a behavior-window study, tokenizer and embedding sweeps, full-eval baselines, an off-policy/on-policy ECPO comparison, and a redistributable MovieLens 1M Colab-T4 run with explicit stage-level parameters.
5. **Failure-aware reproducibility.** We preserve experiment logs for successes, negative results, invalid runs, and post-training collapses, including the conditions under which a GR experiment becomes invalid.

2 Background

2.1 Semantic-ID Generative Recommendation

Semantic-ID GR represents each item as a short sequence of discrete tokens. Recommendation then becomes conditional generation: given a user’s behavior history, a model predicts the token sequence of the next target item. This formulation allows recommendation to reuse autoregressive modeling, beam search, MoE scaling, and post-training alignment ideas from generative AI.

Recent industrial systems have explored this direction from complementary angles. OneRec introduced an end-to-end generative architecture and reported production deployment with scaling

and reinforcement learning evidence [5]. OneRec-V2 improved scalability with a lazy decoder-only design and real-world preference alignment [6]. OneMall unified several e-commerce scenarios with Semantic Tokenization, Transformer-based generation, and RL alignment [4]. GenRec introduced page-wise NTP, token merging, and GRPO-SR for preference-oriented recommendation [7]. OneLive adapted the paradigm to live-streaming with dynamic tokenization and multi-objective alignment [2]. GR4AD studied production advertising with value-aware learning and dynamic beam serving [3].

These papers establish that GR can work in large industrial systems. NANOGENREC addresses a different layer of the problem: the missing reproducible workspace between those system descriptions and a runnable research implementation. It therefore complements prior work by exposing the operational interfaces, experiment governance, and failure records needed to iterate on GR outside the original production stacks.

2.2 Reproducibility Gap

GR papers typically report architecture and production results, while the operational details needed for reproduction are scattered across preprocessing, training, evaluation, and experiment management. Several details are easy to miss:

- Inline evaluation during training is often a health check, not a comparable full-recall metric.
- Feature ablations can become invalid when preprocessing stores a feature but training or inference does not consume it.
- Tokenizer proxy metrics can disagree unless collision, balance, and behavior-aware semantic-neighbor retention are tracked together.
- RL-style post-training can collapse when rollout candidates are sampled from a distribution that no longer matches the current policy.

NANOGENREC treats these details as first-class experimental artifacts.

3 The nanoGenRec Framework

NANOGENREC is organized around one practical goal: a researcher or research agent should be able to move from a paper idea to a runnable experiment, evaluate it with the right contract, and recover the reasoning behind the result weeks later. The framework is therefore centered on research state rather than only model code: hypotheses, paper notes, configs, run queues, validation checks, metrics, failures, and decisions are all first-class artifacts.

3.1 Agentic Research Loop

The repository includes an inbox/outbox protocol, paper notes, decision records, and status files. This is a lightweight workflow for human-agent collaboration over a research line. The agent reads papers, proposes ideas, estimates runtime, writes configs, checks prior experiments, and records decisions. The contribution here is operational: the research process is made inspectable and recoverable rather than embedded in transient chat logs or shell history.

The loop is deliberately failure-aware. Duplicate experiments are surfaced before launch, long jobs are queued instead of chained, inline metrics are separated from reportable full evaluation, and




Figure 1: Agentic research loop implemented by NANOGENREC. The loop turns paper reading, experiment design, execution, evaluation, and decision logging into auditable repository state.

Table 1: Agentic framework artifacts and the research failure modes they address.

Artifact	Research contract	Failure mode addressed
Inbox/outbox	Human instructions and agent findings are explicit	Lost context across sessions
Paper notes	Prior work is summarized into reusable local memory	Re-reading or misremembering papers
YAML configs	Experiments override only intended parameters	Hidden or drifting hyperparameters
Registry hashes	Similar historical runs are surfaced before launch	Duplicate expensive experiments
Queue state	Long jobs are monitored asynchronously	Brittle script chains and forgotten runs
Full-eval reports	Reported numbers use the same evaluation contract	Inline metric overclaiming
Decision records	Negative results and pivots are written down	Cherry-picking and stale conclusions

invalid runs are documented rather than overwritten. This is the main abstraction of the project. Semantic-ID GR is the first demanding workload used to test whether that abstraction can survive a real multi-stage research problem.

Figure 1 shows the agentic control loop. Each transition is backed by a file artifact in the repository rather than by transient chat context. The agent can therefore resume from durable state, while the human collaborator can audit what was proposed, what was run, what failed, and why a direction changed.

Table 1 lists the agent-facing contracts. These artifacts are intentionally mundane: their role is to make research progress recoverable. For example, a YAML config records the exact variant, a registry hash prevents duplicate reruns, the queue state separates long jobs from interactive sessions, and phase README files keep the current best result visible without rereading every log.

3.2 GR Stress-Test Pipeline

The GR validation target consists of six stages:

Table 2: Framework components and the landing checks they make reproducible.

Component	Landing check	Artifact produced
SID tokenizer	Items map to valid multi-token IDs	SID cache and tokenizer metrics
NTP preprocessing	Behavior logs become training/eval shards	Date-bound shard configs
Transformer+MoE NTP	Histories predict next-item SID tokens	Checkpoints and train metadata
Constrained decoding	Generated tokens map back to real items	Full-recall candidate lists
Evaluation	Metrics use generated item IDs	Recall@K reports
Experiment runner	Runs are reproducible and deduplicated	YAML configs and registry hashes
Experiment logs	Results and failures are auditable	EXP notes and phase READMEs

1. **Item representation.** Items are embedded with pretrained representation models. The current experiments use Qwen3 embeddings at several scales.
2. **Semantic ID tokenization.** Embeddings are discretized into 3-token item IDs using residual KMeans and FSQ-family variants.
3. **NTP preprocessing.** User behavior logs are converted into autoregressive training shards, with explicit date windows and evaluation targets.
4. **Generative recommendation.** A Transformer+MoE NTP model predicts item SID tokens from user histories.
5. **Full-recall evaluation.** SID-constrained beam search generates candidate item IDs, and Recall@K is computed with $K \in \{10, 50, 100, 500\}$.
6. **Experiment governance.** YAML configs, registry hashes, duplicate-run checks, queue-based execution, and phase-level README files preserve experiment lineage.

Table 2 summarizes the combined agentic framework and GR validation target as a set of landing checks rather than only software modules.

Table 3 makes the parameter level of each stage explicit. The production-scale empirical line and the public MovieLens line share the same framework interfaces, but they intentionally use different representation and model budgets.

Figure 2 shows the same point as a release diagram: the production line and the public MovieLens line differ in data and representation budget, but meet at the same framework interfaces.

3.3 Experiment Orchestration

New experiments inherit from a shared YAML base config and override only the parameters being tested. Multi-variant experiments are represented as a variant list. A registry hash excludes runtime paths and environment keys, allowing the runner to surface identical or near-identical historical runs before launching a new job. Queue-friendly wrappers support long-running training and evaluation without turning experiment scripts into brittle chains. Each completed run updates a local record of the hypothesis, design, metrics, and interpretation, so the repository accumulates a recoverable experiment lineage rather than only final checkpoints.

Table 3: Stage-level parameterization used by the production empirical study and the public Colab reproduction path.

Stage	Production empirical setting	Public MovieLens 1M setting
Item representation	Qwen3 0.6B / 4B embeddings	Qwen3-Embedding-0.6B over public MovieLens metadata
SID tokenization	3-token RKMeans/FSQ, 4096/8192 codebooks	residual KMeans, 16–64 clusters per layer
Sequence construction	7–90 day windows, up to 445M SID tokens	Public ratings, max 100 items/user, sliding prefixes
NTP model	S/M/L tiers: 17.5M / 71.6M / 101.1M active params	Dense NTP, 2–3 layers, embed_dim=96–128
Training	YAML variants, queued multi-GPU production runs	5–8 epochs, single Colab T4
Alignment	SP-DPO / RF-DPO / GRPO / ECPO	Lightweight GRPO-style exact-SID reward stage
Evaluation	Full SID-constrained recall, R@10/R@500	Beam=1000, 1,000 eval samples, R@1–R@1000
Governance	Duplicate-run checks, EXP logs, phase READMEs	Notebook, metrics, public result page

Production path: Qwen3 item embeddings → RKMeans/FSQ SIDs → NTP shards → Transformer+MoE → preference/RL alignment → constrained full recall → EXP logs

Public path: MovieLens metadata → Qwen3 item embeddings → residual KMeans SIDs → sliding-prefix shards → dense NTP → GRPO-style alignment → constrained full recall → public metrics

Shared interface: item IDs ↔ SID tokens, autoregressive SID sequences, valid-SID decoding, item-level Recall@K, and auditable run artifacts.

Figure 2: Production and public paths share the same framework interfaces while using different data and representation budgets.

3.4 Landing Criteria

We treat GR as landed only when four criteria are met. First, generated token sequences must map back to real item IDs through the tokenizer vocabulary. Second, evaluation must use SID-constrained generation over the item space rather than an inline training proxy. Third, experiment changes must be reproducible from configs and traceable to logs. Fourth, failed runs must be diagnosable from recorded metrics and code-path checks. The empirical study below is organized around these criteria: model scaling tests whether the NTP core works, tokenizer sweeps test whether the item representation is usable, full-recall baselines test whether generation retrieves real items, and post-training experiments test whether alignment remains stable after the SFT stage.

Table 4: NTP model tiers used in the empirical study.

Tier	Embed Dim	Layers	Experts	Active Params
S-tier	256	6	8	$\sim 17.5\text{M}$
M-tier	512	8	8	$\sim 71.6\text{M}$
L-tier	512	12	16	$\sim 101.1\text{M}$

4 Experimental Setup

4.1 Data

The empirical study uses anonymized production behavior data. The largest behavior sweep covers up to 7.85M users and 299.0M raw interactions. After sequence truncation, this corresponds to roughly 445M effective SID tokens. Private data and deployment-specific details are not released; the open-source artifact preserves the code, configs, evaluation protocol, and experiment logs. The main private logs still expose quantitative run scale: for example, the EXP-015 scaling run records 3,042,069 training sequences, 49,575 evaluation targets, 148,533 evaluation positions, 49,383 evaluation items, and 1,000 full-recall samples.

For public reproducibility, the repository also includes a MovieLens path. The strict Colab-T4 run uses public ratings and metadata, Qwen3-Embedding-0.6B item representations, residual KMeans SIDs, 348,363 sliding-prefix NTP training examples, 348,363 alignment examples, a lightweight GRPO-style exact-SID reward stage, constrained decoding, and item-recall evaluation over 1,000 held-out users. An earlier checked-in MovieLens 1M result uses the same users and evaluation split but relies on lightweight hybrid title/genre and co-occurrence features and predates the public reward-alignment stage.

4.2 Compute

The production-scale empirical experiments were run in a real 4 $\times$ NVIDIA H200 GPU environment. This hardware note is included to make the engineering scale of the experiment lineage interpretable; the released code does not depend on that environment. The public MovieLens proof run was executed on a single free Colab T4 GPU.

4.3 Model Tiers

We use three NTP model tiers, summarized in Table 4. Active parameters count only the routed compute used by each token.

4.4 Evaluation Protocol

Full evaluation uses SID-constrained beam search and reports item recall at several cutoffs. Unless otherwise stated, reported R@K values refer to full evaluation rather than inline training evaluation. This distinction matters because inline evaluation uses a limited candidate set and is intended for training health checks.

Table 5: NTP model scaling on a fixed 262M-token dataset.

Config	Active Params	PPL	Loss	R@10	R@500
scale-01	1.7M	235.1	5.460	1.9%	23.6%
scale-02	3.6M	100.4	4.609	3.7%	31.7%
scale-03	5.1M	69.6	4.243	5.4%	45.6%
scale-04	17.5M	28.1	3.334	9.8%	60.5%
scale-05	34.5M	24.0	3.178	11.5%	62.5%
scale-06	71.6M	20.8	3.037	12.6%	66.2%
scale-07	101.1M	19.4	2.965	13.7%	65.8%

5 Framework Validation and Empirical Study

NANOGENREC validates the framework by asking whether the complete loop can reproduce behaviors that matter for GR practitioners. The following sections are therefore not independent benchmark claims. They are evidence that the same framework can support model scaling analysis, data-window analysis, tokenizer diagnosis, full-recall retrieval, and post-training alignment debugging.

5.1 NTP Model Scaling

The first validation target is the NTP core. If the framework has correctly connected SID tokenization, sequence construction, training, and evaluation, model scaling should produce a coherent curve rather than isolated wins. We trained seven model sizes from 1.7M to 101.1M active parameters on a fixed 262M-token dataset built from 31 days of behavior. The corresponding train metadata records 3,042,069 training sequences and 49,575 evaluation targets. Table 5 reports the results.

The fitted loss curve is:

$$L(N) = 2.522 + \frac{2055.1}{N^{0.456}}. \quad (1)$$

The exponent is close to the 0.489 exponent reported by OneRec-V2 [6]. This is an important framework-level signal: the local implementation reproduces a recognizable GR scaling behavior while preserving full-recall metrics. Recall improves substantially through the S-tier and M-tier regime, while gains flatten between 71.6M and 101.1M active parameters under the fixed data budget. Figure 3 shows the fitted curve.

5.2 Behavior-Window Scaling

The second validation target is data plumbing under production time. We fixed the model tier and varied the behavior window. Standard compute-optimal scaling intuition [1] suggests that more data should reduce loss under an i.i.d. assumption. Production recommendation behavior does not satisfy this assumption. Table 6 summarizes the S-tier and M-tier data sweep.

The main observation is non-monotonicity. Extending the window increases the number of users more than the per-user sequence depth, while older behavior can drift away from the current item pool. In this dataset, the exposure window is short enough that old behavior can become stale. Figure 4 visualizes the effect.




Figure 3: NTP scaling law over active parameters.

5.3 Tokenizer and Embedding Sweep

The third validation target is item representation. Tokenizer quality affects the irreducible loss floor and the search space exposed to beam decoding. We evaluated 14 tokenizer variants over 0.6B and 4B embeddings, uniform 4096/8192 codebooks, MGMR-style hierarchical codebooks, and FSQ hidden sizes. The production SID caches used by the embedding-size comparison cover about 1.10M items (1,096,364–1,110,697 items across 0.6B/4B/8B caches), making the tokenizer metrics item-space rather than toy-codebook measurements. Table 7 shows representative results.

The decisive variable in this sweep is the 8192 codebook setting, which improves L2 balance and snHR. The 4B embedding improves semantic-neighbor retention, but also requires enough FSQ capacity to avoid L2 entropy collapse. The current recommended caches are listed in the released experiment logs.

5.4 Full-Recall Baselines

The fourth validation target is the full GR retrieval loop. Table 8 reports current full-eval baselines produced by SID-constrained generation. The strongest observed NTP result is the M-tier model with 4B SID, reaching $R@500=70.4\%$ and $R@10=14.2\%$. Its train metadata records 1,719,646 training sequences, 134.0M training tokens, 49,159 evaluation items, and 1,000 full-recall samples.

5.5 Public MovieLens Reproduction

The private production data cannot be redistributed, so the repository includes a separate public reproduction path on MovieLens. This path is not intended as a public leaderboard benchmark.

Table 6: Behavior-window scaling. The best loss occurs at 14 days, not at the longest window.

Config	Days	Tokens	Users	PPL	R@500
A-7d-S	7	65M	1.02M	30.60	62.1%
B-14d-S	14	130M	1.69M	27.05	58.5%
C-31d-S	31	262M	3.04M	28.05	60.5%
D-62d-S	62	441M	4.86M	30.03	58.6%
E-90d-S	90	553M	6.18M	31.89	56.2%
A-7d-M	7	65M	1.02M	19.31	70.7%
B-14d-M	14	130M	1.69M	18.96	65.8%
C-31d-M	31	262M	3.04M	19.39	65.8%
D-62d-M	62	441M	4.86M	19.80	68.1%

Table 7: Representative tokenizer sweep results. CR is collision rate and snHR is behavior-aware semantic neighbor hit rate.

SID Cache	CR	Gini-d1	Gini-d2	snHR
0.6B nc4096 h64	0.67%	0.3243	0.3546	0.0797
0.6B nc8192 h128	0.42%	0.3914	0.2375	0.0919
4B nc4096 h128	1.60%	0.2877	0.3539	0.1255
4B nc8192 h128	1.28%	0.3192	0.2530	0.1307
4B nc8192x2048 h128	1.41%	0.3191	0.3170	0.1154

Its purpose is to show that the released code can execute the GR framework loop on public data and a widely available free-GPU environment, including Qwen3 item embedding, SID construction, NTP training, reward alignment, and full-recall evaluation in the current notebook.

Table 9 reports the checked-in MovieLens 1M Colab-T4 Qwen+RL reproduction run. The strict public path uses Qwen3 item embeddings, residual KMeans SIDs with a $64 \times 64 \times 64$ token space, a dense 3-layer NTP model with embedding dimension 128, 100 public GRPO-style alignment steps, and SID-constrained beam search with beam size 1000. This table is intended as a run-completion and reproducibility record, not as a public leaderboard comparison.

The strict Qwen+RL run also obtains $R@1=2.2\%$, $R@5=6.7\%$, $R@50=27.9\%$, and target-SID found rate of 88.6%. The result demonstrates that the released artifact can execute the Qwen embedding $\rightarrow$ SID construction $\rightarrow$ NTP training $\rightarrow$ reward alignment $\rightarrow$ constrained full-recall evaluation loop on public data and free-GPU infrastructure. Simple public baselines are kept in the repository for transparency, but the paper uses the MovieLens result only as a reproducibility proof.

5.6 Post-Training Alignment

The final validation target is post-training. We evaluated preference and RL-style alignment methods, including SP-DPO, RF-DPO, GRPO, and ECPO. The most instructive result is the off-policy/on-policy ECPO comparison in Table 10. Both ECPO runs use the same 14-day feature shard with 1,709,380 training sequences and 49,456 evaluation targets; the full evaluation samples 1,000 targets from 26,450 evaluation items.

The off-policy run generates candidates from a reference model whose distribution drifts away




Figure 4: Behavior-window scaling under production temporal drift.

Table 8: Current full-eval baselines.

Config	R@10	R@500	PPL
S-tier bare	11.4%	61.2%	26.52
S-tier with TO-RoPE $t_s=0.5$	11.8%	63.9%	22.7
M-tier bare, 0.6B SID	14.5%	70.2%	18.54
M-tier, 4B SID	14.2%	70.4%	16.55
L-tier with validated options	12.8%	64.1%	20.7

from the current policy. This distorts the importance ratio and causes clipping to dominate. On-policy beam candidates better match the current policy distribution, recovering both PPL and R@500. Figure 5 shows the post-training curves.

6 Discussion

6.1 What the Framework Makes Reproducible

The production data cannot be released. However, the artifact makes the following components inspectable and reusable:

- model and tokenizer implementations;
- preprocessing and evaluation interfaces;
- a CPU-friendly MovieLens smoke path, a strict Colab-T4 Qwen+RL notebook, and a checked-in MovieLens 1M hybrid-feature public run;

Table 9: Public MovieLens 1M Qwen+RL reproduction run.

Run	Users	Items	Train ex.	Align ex.	R@10	R@100	R@500	R@1000
Qwen+RL public path	5,950	3,532	348,363	348,363	10.0%	38.4%	72.2%	86.0%

Table 10: On-policy candidate generation recovers an ECPO collapse.

Config	R@10	R@500	PPL	Clip Rate
Off-policy ECPO	0.7%	2.0%	3791	99%
On-policy ECPO	13.0%	67.8%	14.1	92%
SFT baseline	14.1%	66.2%	16.3	–

- experiment configs and resolved parameter sets;
- full-recall evaluation protocol;
- experiment logs with hypotheses, designs, results, and post-hoc analysis;
- known invalid runs and bug records.

This distinction matters. A GR framework can be useful even when a production dataset is not public, provided that the code path, configuration semantics, evaluation protocol, and empirical record are exposed. The released artifact includes a strict Colab-T4 Qwen+RL notebook plus checked-in MovieLens CPU and hybrid-feature T4 results, demonstrating the public loop at multiple compute levels. The main reproducibility target of NANOGENREC is therefore the working loop: how to construct SIDs, train the model, align it with reward signals, decode valid items, evaluate full recall, organize experiments, and diagnose failures.

6.2 Lessons

Three lessons recur across the experiment lineage. First, landing GR requires full evaluation to be separated from training-time health checks. Second, feature ablations require end-to-end validation across preprocessing, training, and inference; otherwise, experiments can silently compare zeros. Third, post-training alignment in GR is sensitive to candidate distribution, so rollout generation should be treated as part of the algorithm rather than a fixed data-preparation step. These lessons are as much about framework design as model design: most failures arise at component boundaries.

7 Limitations

This report has several limitations. The production behavior data is not released, so external users cannot exactly reproduce the reported production-scale numbers. The repository includes MovieLens CPU and Colab-T4 public paths to verify the released framework loop, but these are reproducibility checks rather than competitive public leaderboard claims. The public GRPO-style stage is intentionally lightweight, and its high clip fraction indicates that alignment hyperparameters remain a tuning target rather than a final recipe. The framework currently focuses on Semantic-ID autoregressive recommendation rather than all forms of generative retrieval. Online serving, latency, and production integration are not the main focus of the artifact. Some post-training methods are preliminary and should be interpreted as framework-validated engineering evidence rather than

Recommendation post-training is an alignment problem, not just a loss-minimization problem




Figure 5: Post-training alignment curves and collapse/recovery behavior.

final algorithmic conclusions. The most important next step is broader public benchmarking across redistributable datasets, so external users can verify not only the code structure but also cross-dataset empirical trends.

8 Conclusion

We presented NANOGENREC, an open-source agentic research framework for landing Semantic-ID generative recommendation. The project contributes a human-agent research protocol, experiment orchestration system, end-to-end GR workflow, full-recall evaluation contract, and 51-experiment lineage. Within this framework, we reproduce several valuable GR outputs: model scaling, non-monotonic behavior-window scaling, tokenizer and embedding effects, full-recall baselines, and post-training alignment collapse/recovery. The central result is not a new standalone algorithm, but a working, inspectable research loop from paper reading and hypothesis formation to Semantic-ID construction, real-item generative retrieval, evaluation, logging, and failure-aware iteration. We hope the artifact helps researchers, practitioners, and research agents turn Generative Recommendation from a collection of system diagrams into an extensible, production-grounded research workflow.

Artifact. The code and experiment logs are available in the public nanoGenRec repository.¹

References

- [1] Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, Tom Hennigan, Eric Noland, Katherine Millican, George van den Driessche, Bogdan Damoc, Aurelia

¹<https://github.com/yzq986/nanoGenRec>

- Guy, Simon Osindero, Karen Simonyan, Erich Elsen, Jack W. Rae, Oriol Vinyals, and Laurent Sifre. Training compute-optimal large language models. In *Advances in Neural Information Processing Systems*, 2022. URL <https://arxiv.org/abs/2203.15556>.
- [2] Shen Wang, Yusheng Huang, Ruochen Yang, et al. Onelive: Dynamically unified generative framework for live-streaming recommendation, 2026. URL <https://arxiv.org/abs/2602.08612>.
- [3] Ben Xue, Dan Liu, Lixiang Wang, et al. Generative recommendation for large-scale advertising, 2026. URL <https://arxiv.org/abs/2602.22732>.
- [4] Kun Zhang, Jingming Zhang, Wei Cheng, et al. Onemall: One architecture, more scenarios – end-to-end generative recommender family at kuaishou e-commerce, 2026. URL <https://arxiv.org/abs/2601.21770>.
- [5] Guorui Zhou et al. Onerec technical report, 2025. URL <https://arxiv.org/abs/2506.13695>.
- [6] Guorui Zhou et al. Onerec-v2 technical report, 2025. URL <https://arxiv.org/abs/2508.20900>.
- [7] Yanyan Zou, Junbo Qi, Lunsong Huang, Yu Li, Kewei Xu, Jiahao Gao, Binglei Zhao, Xuanhua Yang, Sulong Xu, and Shengjie Li. Genrec: A preference-oriented generative framework for large-scale recommendation, 2026. URL <https://arxiv.org/abs/2604.14878>.